



School of Engineering

Introduction Programming Massively Parallel Processors

Dr. Maria Pantoja



**Santa Clara
University**

School of Engineering

Motivation

Text & Notes

References:

- NVIDIA. *The NVIDIA CUDA Programming Guide*.
- NVIDIA. CUDA Reference Manual.
- CUDA by Example An Introduction to General-Purpose GPU programming. 2010
- Kirk & Hwu. *Programming Massively Parallel Processors: A Hands-on Approach*. 2012.


Moore's Law (paraphrased)

“The number of transistors on an integrated circuit doubles every two years.”

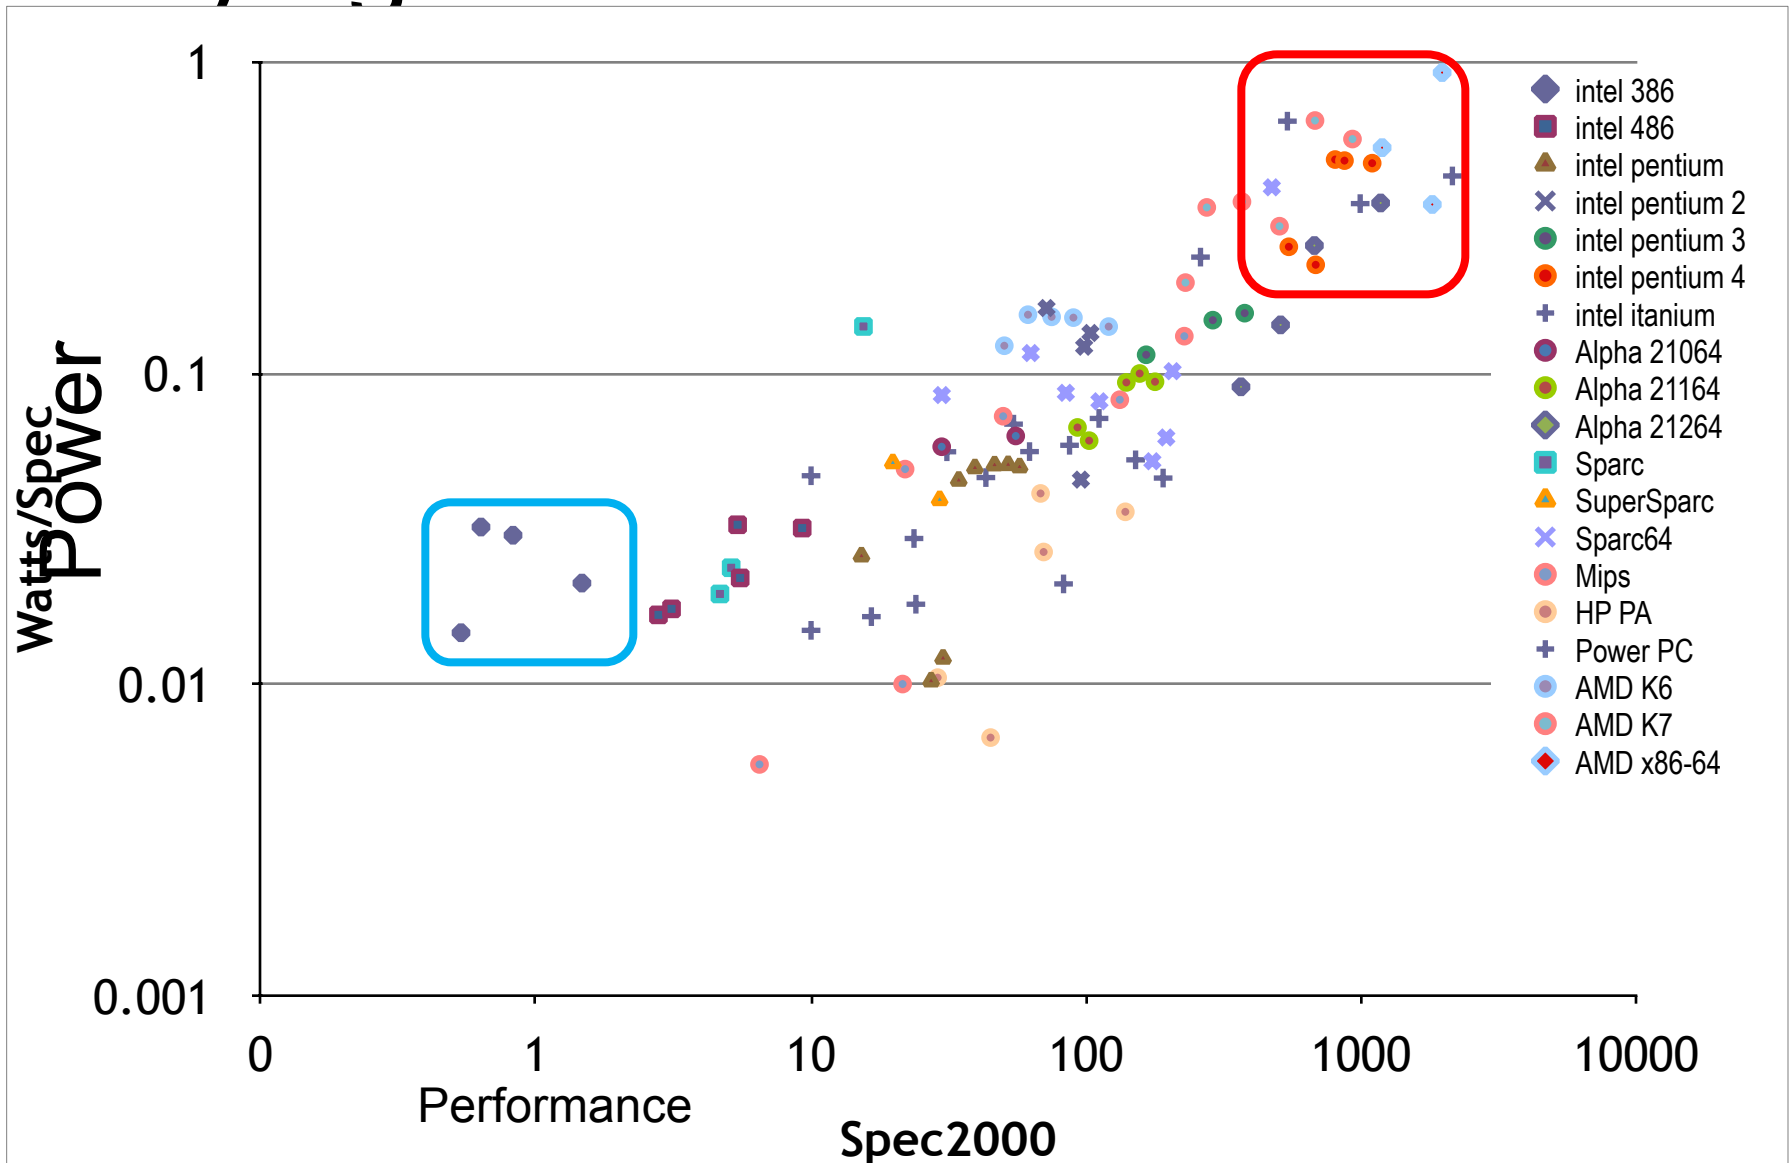
– Gordon E. Moore

Moore's Law (Visualized)

CPU Transistor Counts 1971-2008 & Moore's Law



Buying Performance with Power



(courtesy Mark Horowitz and Kevin Skadron)

Serial Performance Scaling is Over


- **Cannot** continue to scale processor frequencies
 - no 10 GHz chips
- **Cannot** continue to increase power consumption
 - can't melt chip
- Can continue to increase transistor density
 - as per Moore's Law

How to Use Transistors?

- Instruction-level parallelism
 - out-of-order execution, speculation, ...
 - **vanishing opportunities** in power-constrained world
- Data-level parallelism
 - vector units, SIMD execution, ...
 - increasing ... SSE, AVX, Cell SPE, Clearspeed, GPU
- Thread-level parallelism
 - increasing ... multithreading, multicore, manycore
 - Intel Core2, AMD Phenom, Sun Niagara, STI Cell, NVIDIA Fermi, ...

Why Massively Parallel Processing?

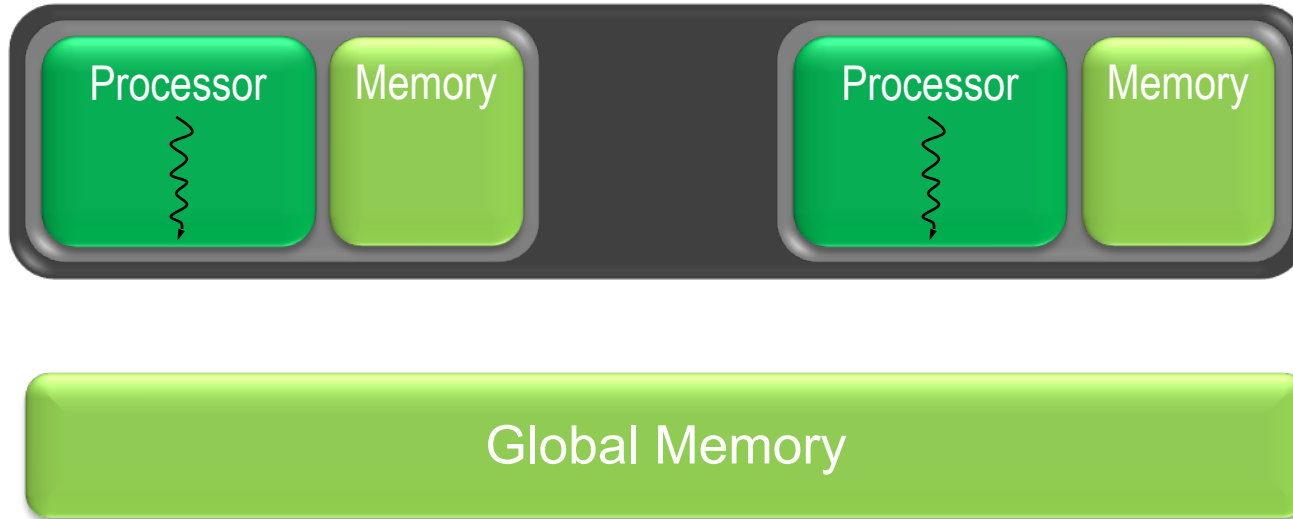
- A quiet revolution and potential build-up
 - Computation: TFLOPs vs. 100 GFLOPs



The “New” Moore’s Law

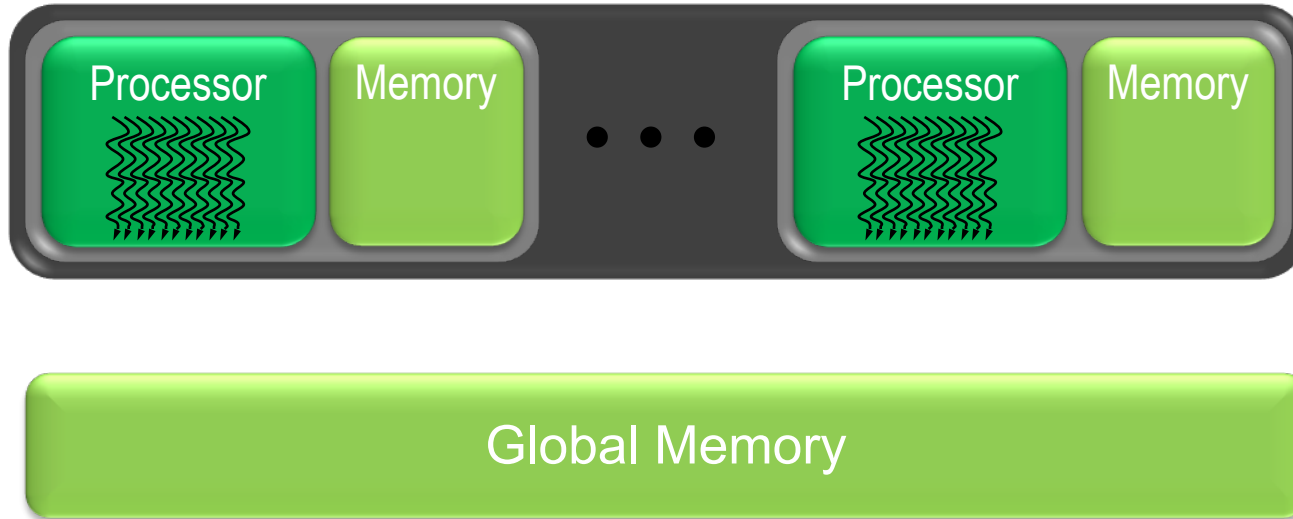
- Computers no longer get faster, just wider
- You *must* re-think your algorithms to be parallel !
- Data-parallel computing is most scalable solution
 - Otherwise: refactor code for 2 cores
 - You will always have more data than cores – build the computation around the data

Generic Multicore Chip



- Handful of processors each supporting ~1 hardware thread
- On-chip memory near processors (cache, RAM, or both)
- Shared global memory space (external DRAM)

Generic Manycore Chip



- Many processors each supporting **many hardware threads**
- **On-chip memory** near processors (cache, RAM, or both)
- **Shared global memory** space (external DRAM)

Enter the GPU

- Massive economies of scale
- Massively parallel




Why is this different from a CPU?

- Different goals produce different designs
 - GPU assumes work load is highly parallel
 - CPU must be good at everything, parallel or not
- CPU: **minimize latency** experienced by 1 thread
 - big on-chip caches
 - sophisticated control logic
- GPU: **maximize throughput** of all threads
 - # threads in flight limited by resources => lots of resources (registers, bandwidth, etc.)
 - multithreading can hide latency => skip the big caches
 - share control logic across many threads

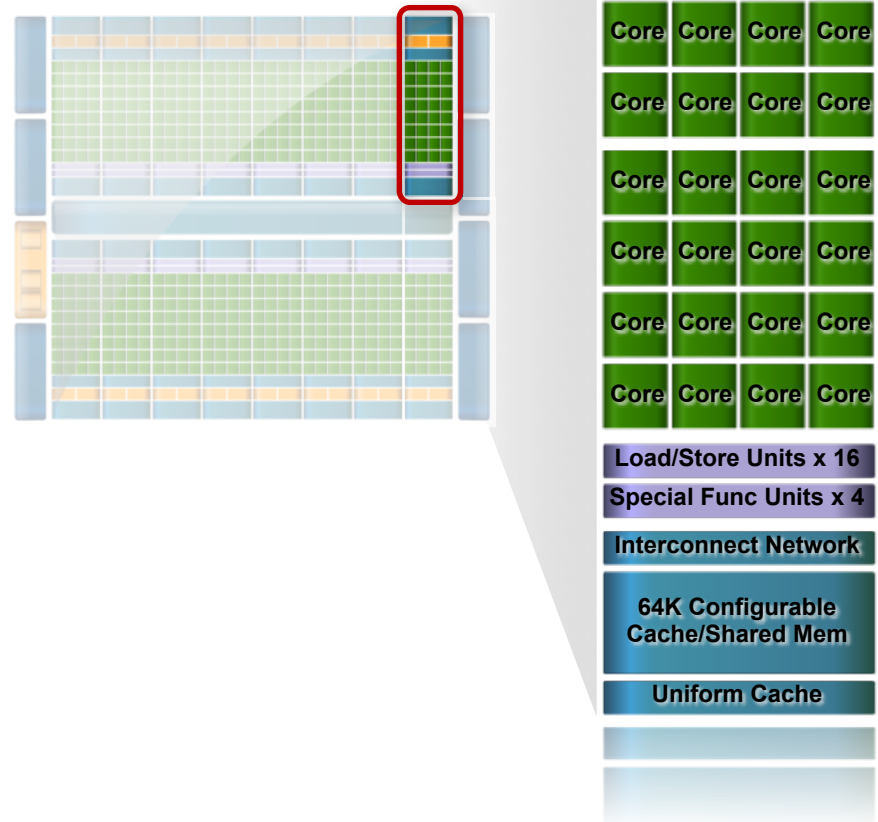
NVIDIA GPU Architecture

Fermi GF100



SM Multiprocessor

- 32 CUDA Cores per SM (512 total)
- Direct load/store to memory
 - Usual linear sequence of bytes
 - High bandwidth (Hundreds GB/sec)
- 64KB of fast, on-chip RAM
 - Software or hardware-managed
 - Shared amongst CUDA cores
 - Enables thread communication



Key Architectural Ideas

- SIMT (Single Instruction Multiple Thread) execution
 - threads run in groups of 32 called warps
 - threads in a warp share instruction unit (IU)
 - HW automatically handles divergence
- Hardware multithreading
 - HW resource allocation & thread scheduling
 - HW relies on threads to hide latency
- Threads have all resources needed to run
 - any warp not waiting for something can run
 - context switching is (basically) free



NVIDIA GPU Architecture

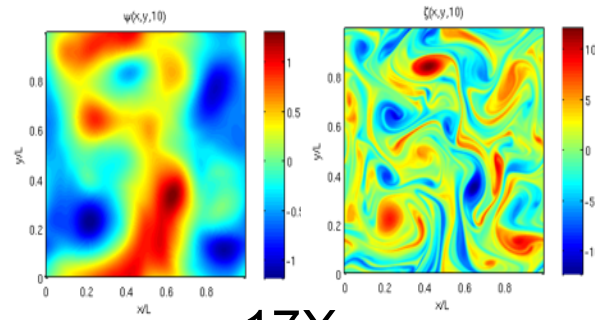
Kepler

CUDA


- Scalable parallel programming model
- Minimal extensions to familiar C/C++ environment
- Heterogeneous serial-parallel computing




45X



17X




100X




13–457x

Motivation



110-240X

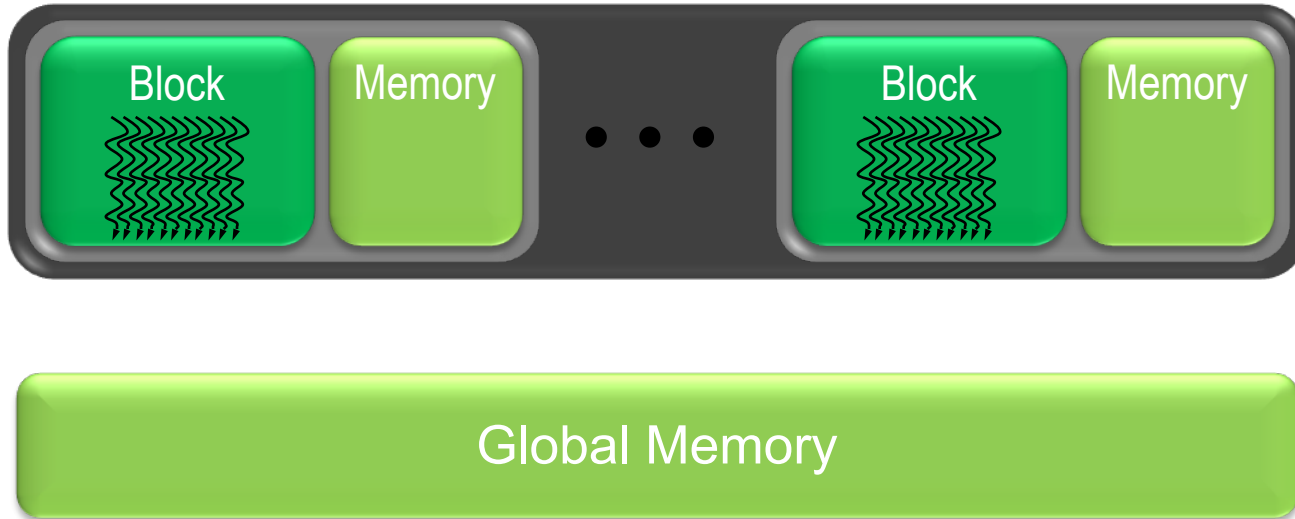


35X


C_{UDA}: Scalable parallel programming

- Augment C/C++ with minimalist abstractions
 - let programmers focus on parallel algorithms
 - *not* mechanics of a parallel programming language
- Provide straightforward mapping onto hardware
 - good fit to GPU architecture
 - maps well to multi-core CPUs too
- Scale to 100s of cores & 10,000s of parallel threads
 - GPU threads are lightweight — create / switch is free
 - GPU needs 1000s of threads for full utilization

CUDA Model of Parallelism



Speedup of Applications



- GeForce 8800 GTX vs. 2.2GHz Opteron 248
- 10× speedup in a kernel is typical, as long as the kernel can occupy enough parallel threads
- 25× to 400× speedup if the function's data requirements and control flow suit the GPU and the application is optimized

Final Thoughts

- Parallel hardware is here to stay
- GPUs are massively parallel manycore processors
 - easily available and fully programmable
- Parallelism & scalability are crucial for success
- This presents many important research challenges
 - not to speak of the educational challenges

Course Equipment

- Your own PCs with a CUDA-enabled GPU
- Design Center

- NVIDIA GeForce GTX 570 boards

Fermi Architecture GPUs

- Little Fe cluster
 - Alcatraz1 and 2 with

Kepler

- K20 comp cap 3.5
 - GTX 690 comp cap 3.0

Fermi

GTX 480

Course Equipment

- Design Center

- Remote Login Instructions:

- <http://www.scu.edu/engineering/centers/scudc/Terminal-Services.cfm>

- Using putty

- <http://www.cse.scu.edu/~sfigueira/10/remote.html>

- (linux.scudc.scu.edu)

- Login to alcatraz 2:

Your Own Computer

- Nvidia CUDA getting started guide for MS Windows
 - CUDA enabled GPU
 - MS Windows XP, vista , or 7 (better 7)
 - Device driver
 - Cuda Software (nvcc compiler)
 - MS Visual Studio Express edition (free for students) 2005 or later (better 2010)

Your Own Computer

- Verify you have a CUDA enabled GPU
- List of CUDA enabled GPUs:

http://www.nvidia.com/object/cuda_gpus.html

Your Own Computer

- Download CUDA software:
 - Driver.
 - CUDA toolkit:

Tools needed to compile and build CUDA applications
 - SDK

Sample projects that have all necessary project configuration and build in files to perform one click builds using MS Visual Studio

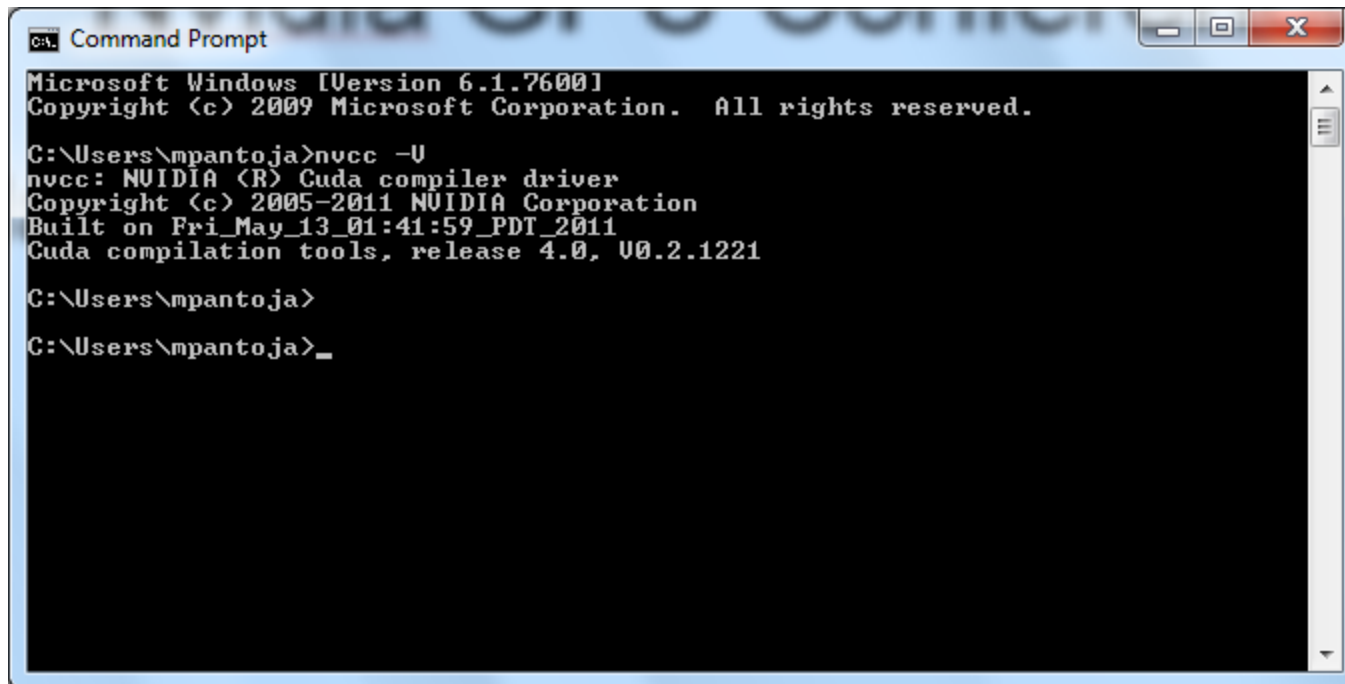
<http://developer.nvidia.com/cuda-toolkit-41>

Your Own Computer

- Verify the installation

Command prompt window:

Start->all programs->accessories->command prompt



```
Command Prompt
Microsoft Windows [Version 6.1.7600]
Copyright (c) 2009 Microsoft Corporation. All rights reserved.

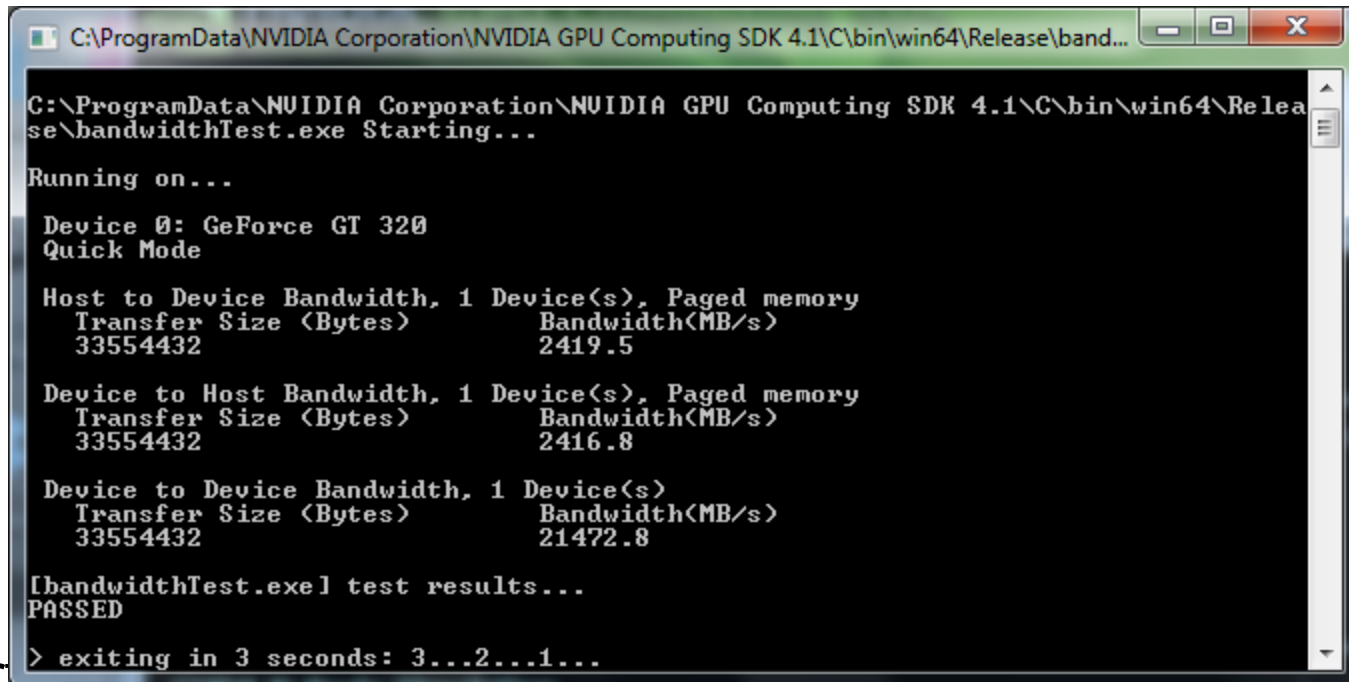
C:\Users\mpantoja>nvcc -V
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2011 NVIDIA Corporation
Built on Fri_May_13_01:41:59_PDT_2011
Cuda compilation tools, release 4.0, V0.2.1221

C:\Users\mpantoja>
C:\Users\mpantoja>_
```

Your Own Computer

- Verify the installation

Bandwidthtest: sdk\C\src\bandwidthtest



The screenshot shows a Windows command prompt window titled "C:\ProgramData\NVIDIA Corporation\NVIDIA GPU Computing SDK 4.1\C\bin\win64\Release\band...". The window contains the following text:

```
C:\ProgramData\NVIDIA Corporation\NVIDIA GPU Computing SDK 4.1\C\bin\win64\Release\bandwidthTest.exe Starting...
Running on...
Device 0: GeForce GT 320
Quick Mode

Host to Device Bandwidth, 1 Device(s), Paged memory
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                   2419.5

Device to Host Bandwidth, 1 Device(s), Paged memory
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                   2416.8

Device to Device Bandwidth, 1 Device(s)
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                   21472.8

[bandwidthTest.exe] test results...
PASSED

> exiting in 3 seconds: 3...2...1...
```

Run part of the test to verify the installation is correct

Cuda and MSVC++ 2010

1. Open MS VC++ 2010.

2. Create new project:

File->new->project, select empty project

3. To get cuda syntax highlighting:

tools->options->text editor->file extensions

type *.cu in input box and click on apply

4. Right click on project, select build Customizations, click on the cuda 4.1 select box

5. Right click on project->properties->configuration properties->linker->input

add cudart.lib to the list of additional dependencies

6. Add new item to project->C++ file but name the file with *.cu extension

Your Own Computer

Compile new CUDA Project

`sdk\C\src\bandwidthtest`

Solution should be placed on debugging folder